


(a) World-model training convergence



(b) Train vs held-out generalization

